

Classification Trees in R

May 15, 2016

Contents

Grow one tree	2
Bootstrap aggregating (bagging)	8
Boosting	10
Random Forest	13
Making prettier trees	17

We'll be working with the same Twitter dataset again this week:

```
library(dplyr)
library(ggplot2)
library(scales)
library(caret)

twitter = read.delim('bot_or_not.tsv',
                    sep = '\t',
                    header = TRUE)
```

As usual, divide the data into test and train.

```
# tell R which variables are categorical (factors)
twitter$bot = factor(twitter$bot)
twitter$default_profile = factor(twitter$default_profile)
twitter$default_profile_image = factor(twitter$default_profile_image)
twitter$geo_enabled = factor(twitter$geo_enabled)

summary(twitter)
```



```
##  bot      statuses_count  default_profile default_profile_image
##  0:2672  Min.   :    0    0:2256          0:3077
##  1: 504   1st Qu.:  188   1: 920          1: 99
##                Median :   723
##                Mean    :  3277
##                3rd Qu.:  2646
##                Max.    :137264
##  friends_count  followers_count  favourites_count geo_enabled
##  Min.   :    11  Min.   :    0.0  Min.   :    0  0:1773
##  1st Qu.:   300  1st Qu.:   95.0  1st Qu.:   14  1:1403
##  Median :   615  Median :  288.0  Median :   122
##  Mean   :  2358  Mean   : 3709.3  Mean   :  1100
##  3rd Qu.:  1229  3rd Qu.:  830.5  3rd Qu.:   593
```

```
## Max. :1175187 Max. :1396699.0 Max. :176219
## listed_count account_age_hours diversity
## Min. : 0.00 Min. : 2072 Min. :0.0050
## 1st Qu.: 4.00 1st Qu.:30285 1st Qu.:0.6254
## Median : 16.00 Median :47484 Median :0.6963
## Mean : 84.77 Mean :43664 Mean :0.6791
## 3rd Qu.: 51.00 3rd Qu.:56718 3rd Qu.:0.7626
## Max. :9491.00 Max. :78841 Max. :1.0000
## mean_mins_between_tweets mean_tweet_length mean_retweets
## Min. : -15.7 Min. : 8.50 Min. : 1.000
## 1st Qu.: 1152.8 1st Qu.: 80.79 1st Qu.: 1.167
## Median : 3851.7 Median : 91.74 Median : 1.636
## Mean : 14715.4 Mean : 91.41 Mean : 3.873
## 3rd Qu.: 10823.8 3rd Qu.:103.28 3rd Qu.: 2.424
## Max. :1139015.0 Max. :287.88 Max. :1961.300
## reply_rate
## Min. :0.0000
## 1st Qu.:0.1232
## Median :0.3137
## Mean :0.3411
## 3rd Qu.:0.5279
## Max. :1.0000
```

```
set.seed(243)

# select the training observations
in_train = createDataPartition(y = twitter$bot,
                                p = 0.75, # 75% in train, 25% in test
                                list = FALSE)

training_set = twitter[in_train, ]
testing_set = twitter[-in_train, ]
```

Grow one tree

caret has lots of different tree models, so check 'em out. We can make a simple tree model using the `rpart` method.

```
tree_model = train(bot ~.,
                    method = 'rpart',
                    data = training_set)

print(tree_model)
```

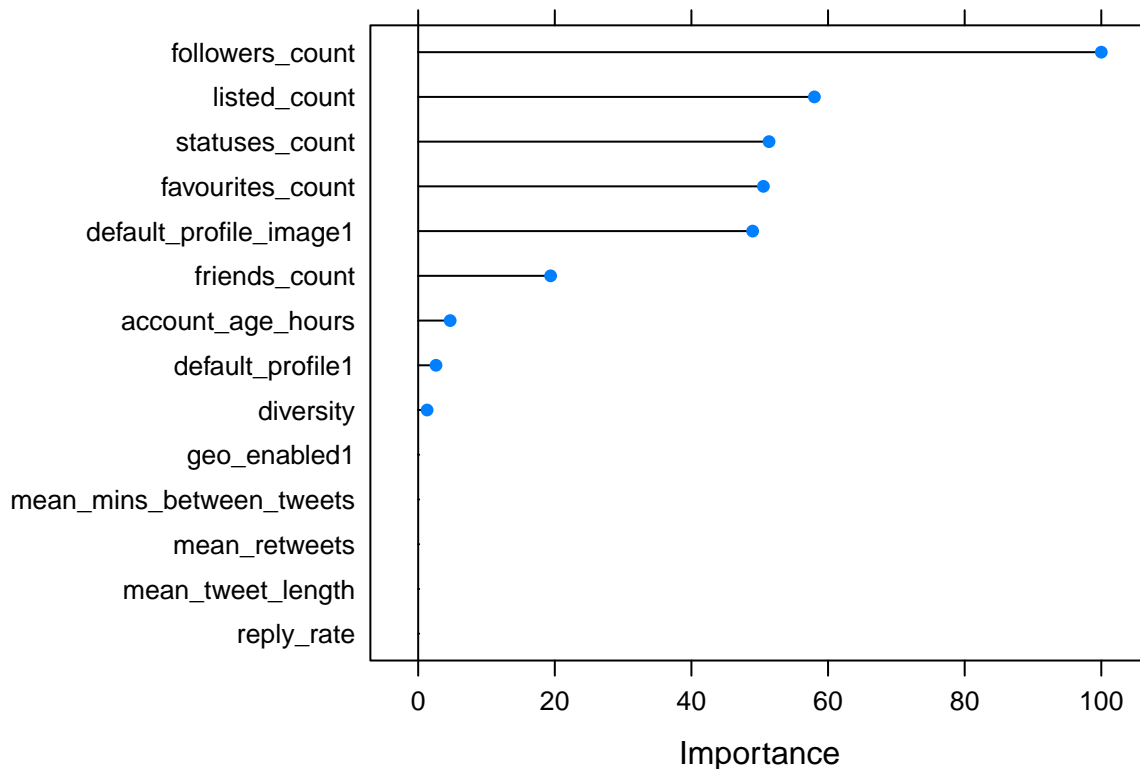
```
## CART
##
## 2382 samples
## 14 predictor
## 2 classes: '0', '1'
##
## No pre-processing
```

```
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2382, 2382, 2382, 2382, 2382, 2382, ...
## Resampling results across tuning parameters:
##
##      cp          Accuracy      Kappa
##  0.01851852  0.9060279  0.5984621
##  0.03306878  0.9006168  0.5525340
##  0.36507937  0.8652871  0.2341893
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.01851852.
```


```
print(tree_model$finalModel)
```

```
## n= 2382
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 2382 378 0 (0.84130982 0.15869018)
##    2) followers_count>=31.5 2136 186 0 (0.91292135 0.08707865) *
##    3) followers_count< 31.5 246  54 1 (0.21951220 0.78048780)
##      6) friends_count< 286.5 96  44 0 (0.54166667 0.45833333)
##        12) friends_count>=102 37  6 0 (0.83783784 0.16216216) *
##        13) friends_count< 102 59 21 1 (0.35593220 0.64406780) *
##          7) friends_count>=286.5 150  2 1 (0.01333333 0.98666667) *
```

```
plot(varImp(tree_model))
```



```
# plot the tree!
plot(tree_model$finalModel)
text(tree_model$finalModel, use.n = TRUE, all = TRUE, cex = 0.60)
```



```
# test the predictions
tree_predictions = predict(tree_model, newdata = testing_set)
confusionMatrix(tree_predictions, testing_set$bot)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 656  72
##           1  12  54
##
##           Accuracy : 0.8942
##           95% CI : (0.8707, 0.9147)
##           No Information Rate : 0.8413
##           P-Value [Acc > NIR] : 1.162e-05
##
##           Kappa : 0.5089
##           McNemar's Test P-Value : 1.215e-10
##
##           Sensitivity : 0.9820
##           Specificity : 0.4286
##           Pos Pred Value : 0.9011
##           Neg Pred Value : 0.8182
##           Prevalence : 0.8413
##           Detection Rate : 0.8262
##           Detection Prevalence : 0.9169
##           Balanced Accuracy : 0.7053
##
##           'Positive' Class : 0
##
```

By default, the train function will try three values of the complexity parameter, but we can tell it to try more using the `tuneLength` argument.

```
tree_model = train(bot ~.,
                    method = 'rpart',
                    data = training_set,
                    tuneLength = 10)
print(tree_model)
```


```
## CART
##
## 2382 samples
## 14 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2382, 2382, 2382, 2382, 2382, 2382, ...
## Resampling results across tuning parameters:
##
##   cp          Accuracy   Kappa
## 0.002645503 0.8960147 0.5782026
## 0.003968254 0.8975288 0.5804829
## 0.005291005 0.8991581 0.5835626
## 0.006613757 0.8992922 0.5821151
## 0.009259259 0.8997177 0.5782421
## 0.015873016 0.8979459 0.5597171
## 0.016402116 0.8980368 0.5597664
## 0.018518519 0.8978292 0.5529883
## 0.033068783 0.8952924 0.5260809
## 0.365079365 0.8801403 0.3894097
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.009259259.
```

```
print(tree_model$finalModel)
```


```
## n= 2382
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 2382 378 0 (0.84130982 0.15869018)
## 2) followers_count>=31.5 2136 186 0 (0.91292135 0.08707865)
## 4) friends_count>=99.5 2068 154 0 (0.92553191 0.07446809)
## 8) followers_count>=150.5 1540 67 0 (0.95649351 0.04350649) *
## 9) followers_count< 150.5 528 87 0 (0.83522727 0.16477273)
## 18) friends_count< 529 418 29 0 (0.93062201 0.06937799) *
## 19) friends_count>=529 110 52 1 (0.47272727 0.52727273)
## 38) followers_count>=75 83 34 0 (0.59036145 0.40963855)
## 76) friends_count>=968 25 4 0 (0.84000000 0.16000000) *
## 77) friends_count< 968 58 28 1 (0.48275862 0.51724138)
## 154) statuses_count>=335 20 5 0 (0.75000000 0.25000000) *
## 155) statuses_count< 335 38 13 1 (0.34210526 0.65789474) *
## 39) followers_count< 75 27 3 1 (0.11111111 0.88888889) *
```

```
##      5) friends_count< 99.5 68  32 0 (0.52941176 0.47058824)
##      10) followers_count>=73 34  10 0 (0.70588235 0.29411765) *
##      11) followers_count< 73 34  12 1 (0.35294118 0.64705882) *
##      3) followers_count< 31.5 246  54 1 (0.21951220 0.78048780)
##      6) friends_count< 286.5 96  44 0 (0.54166667 0.45833333)
##      12) friends_count>=102 37   6 0 (0.83783784 0.16216216) *
##      13) friends_count< 102 59  21 1 (0.35593220 0.64406780)
##      26) friends_count< 21.5  7   0 0 (1.00000000 0.00000000) *
##      27) friends_count>=21.5 52  14 1 (0.26923077 0.73076923) *
##      7) friends_count>=286.5 150   2 1 (0.01333333 0.98666667) *
```

```
plot(varImp(tree_model))
```



```
# plot accuracy by the complexity parameter
plot(tree_model)
```



```
# test the predictions
```

```
tree_predictions = predict(tree_model, newdata = testing_set)
confusionMatrix(tree_predictions, testing_set$bot)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 639  52
```

```
##           1  29  74
```

```
##
```

```
##           Accuracy : 0.898
```

```
##           95% CI : (0.8748, 0.9182)
```

```
## No Information Rate : 0.8413
```

```
## P-Value [Acc > NIR] : 2.585e-06
```

```
##
```

```
##           Kappa : 0.5874
```

```
## McNemar's Test P-Value : 0.01451
```

```
##
```

```
##           Sensitivity : 0.9566
```

```
##           Specificity : 0.5873
```

```
## Pos Pred Value : 0.9247
```

```
## Neg Pred Value : 0.7184
```

```
## Prevalence : 0.8413
```

```
## Detection Rate : 0.8048
```

```
## Detection Prevalence : 0.8703
```

```
## Balanced Accuracy : 0.7719
```

```
##
```

```
##          'Positive' Class : 0
##
```

Bootstrap aggregating (bagging)


You might have to install some extra packages before this one will run. The key idea in bagging is that we resample the input data and recompute the predictions. Then, use the average or majority vote to determine the class.

```
bagged_model = train(bot ~.,
                     method = 'treebag',
                     data = training_set)

print(bagged_model)
```

```
## Bagged CART
##
## 2382 samples
##   14 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2382, 2382, 2382, 2382, 2382, 2382, ...
## Resampling results:
##
##   Accuracy   Kappa
##   0.9099259  0.6238812
##
##
```

```
plot(varImp(bagged_model))
```

```
bagged_predictions = predict(bagged_model, testing_set)
confusionMatrix(bagged_predictions, testing_set$bot)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 652  51
##           1  16  75
##
##           Accuracy : 0.9156
##           95% CI : (0.8941, 0.934)
##           No Information Rate : 0.8413
##           P-Value [Acc > NIR] : 4.425e-10
##
##           Kappa : 0.6438
##           McNemar's Test P-Value : 3.271e-05
##
##           Sensitivity : 0.9760
##           Specificity : 0.5952
##           Pos Pred Value : 0.9275
##           Neg Pred Value : 0.8242
##           Prevalence : 0.8413
##           Detection Rate : 0.8212
##           Detection Prevalence : 0.8854
##           Balanced Accuracy : 0.7856
##
##           'Positive' Class : 0
##
```

In this case, we do get some accuracy gains from bagging.

Boosting


The key idea of boosting is that we amplify the signal of weak predictors by up-weighting misclassified observations at each split point.

```
boost_model = train(bot ~.,
                    method = 'gbm',
                    data = training_set,
                    verbose = FALSE)


print(boost_model)
```

```
## Stochastic Gradient Boosting
##
## 2382 samples
## 14 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2382, 2382, 2382, 2382, 2382, 2382, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  Accuracy  Kappa
##  1                  50      0.9026786  0.5660969
##  1                  100     0.9083585  0.5998731
##  1                  150     0.9117005  0.6194439
##  2                   50     0.9119541  0.6166607
##  2                  100     0.9172078  0.6519053
##  2                  150     0.9180795  0.6591229
##  3                   50     0.9153706  0.6402409
##  3                  100     0.9188683  0.6638209
##  3                  150     0.9179840  0.6639106
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 100,
##  interaction.depth = 3, shrinkage = 0.1 and n.minobsinnode = 10.
```


```
plot(boost_model)
```



```
plot(varImp(boost_model))
```



```
#TODO: remove this?
summary(boost_model$finalModel)
```



```
##                                var    rel.inf
## followers_count               followers_count 47.2144335
```

```
## friends_count                friends_count 29.0930492
## default_profile1             default_profile1 4.6327624
## statuses_count               statuses_count 3.8931255
## account_age_hours            account_age_hours 3.4048035
## favourites_count              favourites_count 2.5434848
## geo_enabled1                 geo_enabled1 1.7800263
## diversity                     diversity 1.6530707
## default_profile_image1        default_profile_image1 1.6402545
## mean_tweet_length             mean_tweet_length 1.1121697
## mean_mins_between_tweets      mean_mins_between_tweets 1.1080210
## reply_rate                    reply_rate 0.8990138
## listed_count                  listed_count 0.6135349
## mean_retweets                 mean_retweets 0.4122503
```

```
# predict
boost_predictions = predict(boost_model, testing_set)
confusionMatrix(boost_predictions, testing_set$bot)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 659  49
##              1   9  77
##
##              Accuracy : 0.927
##              95% CI : (0.9066, 0.9441)
##              No Information Rate : 0.8413
##              P-Value [Acc > NIR] : 3.324e-13
##
##              Kappa : 0.686
##              McNemar's Test P-Value : 3.040e-07
##
##              Sensitivity : 0.9865
##              Specificity : 0.6111
##              Pos Pred Value : 0.9308
##              Neg Pred Value : 0.8953
##              Prevalence : 0.8413
##              Detection Rate : 0.8300
##              Detection Prevalence : 0.8917
##              Balanced Accuracy : 0.7988
##
##              'Positive' Class : 0
##
```

Random Forest


Random forest is a bagging method where we resample both observations, and variables, grow multiple trees and aggregate votes. It's one of the most accurate classifiers, but can be slow. Might want to run this one at home...

```
rf_model = train(bot ~.,
                 data = training_set,
                 method = 'rf',
                 prox = TRUE,
                 verbose = TRUE)
```


```
print(rf_model)
```

```
## Random Forest
##
## 2382 samples
## 14 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2382, 2382, 2382, 2382, 2382, 2382, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9078537 0.5868531
## 8 0.9122775 0.6217937
## 14 0.9118290 0.6279741
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 8.
```

```
plot(rf_model)
```



```
plot(rf_model$finalModel)
```



```
# pull a tree out of the forest
```

```
head(getTree(rf_model$finalModel, k = 5, labelVar = TRUE))
```

```
##   left daughter right daughter   split var split point status
## 1         2         3   listed_count      3.5         1
## 2         4         5  friends_count    529.0         1
## 3         6         7  friends_count     99.5         1
## 4         8         9  friends_count    100.5         1
## 5        10        11 followers_count     77.5         1
## 6        12        13   listed_count     18.0         1
## prediction
## 1      <NA>
## 2      <NA>
## 3      <NA>
## 4      <NA>
## 5      <NA>
## 6      <NA>
```

```
# predict
```

```
rf_predictions = predict(rf_model, testing_set)
confusionMatrix(rf_predictions, testing_set$bot)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    0    1
##              0 656  55
##              1  12  71
##
##              Accuracy : 0.9156
##              95% CI : (0.8941, 0.934)
##      No Information Rate : 0.8413
##      P-Value [Acc > NIR] : 4.425e-10
##
##              Kappa : 0.6332
##  McNemar's Test P-Value : 2.880e-07
##
##      Sensitivity : 0.9820
##      Specificity : 0.5635
##      Pos Pred Value : 0.9226
##      Neg Pred Value : 0.8554
##      Prevalence : 0.8413
##      Detection Rate : 0.8262
##      Detection Prevalence : 0.8955
##      Balanced Accuracy : 0.7728
##
##      'Positive' Class : 0
##
```

As always, we can compare the models with the `resamples` function.

```
# compare the three methods
results = resamples(list(tree_model = tree_model,
                        bagged_model = bagged_model,
                        boost_model = boost_model,
                        rf_model = rf_model))


# compare accuracy and kappa
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: tree_model, bagged_model, boost_model, rf_model
## Number of resamples: 25
##
## Accuracy
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## tree_model  0.8781  0.8916 0.8996 0.8997  0.9061 0.9195    0
## bagged_model 0.8920  0.9051 0.9121 0.9099  0.9156 0.9240    0
## boost_model  0.9082  0.9128 0.9194 0.9189  0.9234 0.9352    0
## rf_model     0.8954  0.9051 0.9121 0.9123  0.9176 0.9308    0
##
## Kappa
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## tree_model  0.5085  0.5502 0.5750 0.5782  0.6050 0.6528    0
## bagged_model 0.5717  0.5988 0.6288 0.6239  0.6496 0.6756    0
```



```
## boost_model  0.6110  0.6432 0.6634 0.6638  0.6825 0.7134    0
## rf_model     0.5548  0.5936 0.6190 0.6218  0.6498 0.6817    0
```

```
# plot results
dotplot(results)
```



How do the tree models compare with logistic regression?

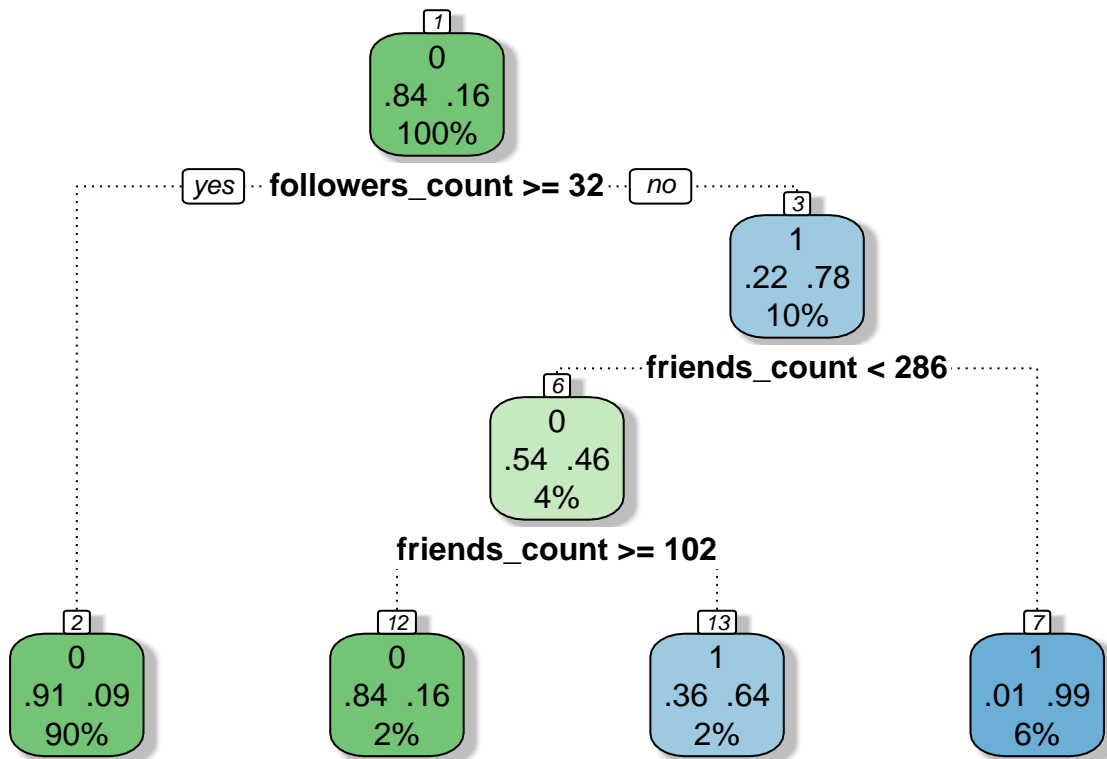
Making prettier trees

If you want to make your trees look a little prettier, try out the `rattle` package.

```
library(rattle)

tree_model = train(bot ~.,
                   method = 'rpart',
                   data = training_set)

# plot the final model
fancyRpartPlot(tree_model$finalModel)
```



Rattle 2016-May-16 11:22:47 erin