

Logistic regression in R

May 09, 2016

Contents

Introduction	1
Exploratory data analysis	1
Feature engineering	12
Logistic Regression	15
Training and testing sets	15
Training logistic regressions	15

Introduction

For the second project we'll explore user data from Twitter to identify accounts likely belonging to bots. The data set has variables about profile configuration (`default_profile`, `default_profile_image`), connectivity (`friends_count`, `followers_count`), and some information about the nature of their tweets (`diversity`, `mean_mins_between_tweets`). Additionally, there's an outcome variable called `bot` that denotes whether the account belongs to a bot (`bot == 1`) or to a human (`bot == 0`).

```
library(dplyr)
library(ggplot2)
library(GGally)
library(caret)

twitter = read.delim('bot_or_not.tsv',
                    sep = '\t',
                    header = TRUE)
```

Exploratory data analysis

We've got a brand new data set, so let's familiarize ourselves by conducting an exploratory data analysis. Let's start by summarizing the whole data set to see what the variable values are.

```
summary(twitter)
```

##	bot	statuses_count	default_profile	default_profile_image
##	Min. :0.0000	Min. : 0	Min. :0.0000	Min. :0.00000
##	1st Qu.:0.0000	1st Qu.: 188	1st Qu.:0.0000	1st Qu.:0.00000
##	Median :0.0000	Median : 723	Median :0.0000	Median :0.00000
##	Mean :0.1587	Mean : 3277	Mean :0.2897	Mean :0.03117
##	3rd Qu.:0.0000	3rd Qu.: 2646	3rd Qu.:1.0000	3rd Qu.:0.00000
##	Max. :1.0000	Max. :137264	Max. :1.0000	Max. :1.00000

```
## friends_count      followers_count      favourites_count  geo_enabled
## Min.      :    11  Min.      :    0.0  Min.      :    0  Min.      :0.0000
## 1st Qu.:    300  1st Qu.:    95.0  1st Qu.:    14  1st Qu.:0.0000
## Median :    615  Median :   288.0  Median :    122  Median :0.0000
## Mean      :   2358  Mean      :  3709.3  Mean      :   1100  Mean      :0.4418
## 3rd Qu.:   1229  3rd Qu.:   830.5  3rd Qu.:    593  3rd Qu.:1.0000
## Max.      :1175187  Max.      :1396699.0  Max.      :176219  Max.      :1.0000
## listed_count      account_age_hours      diversity
## Min.      :    0.00  Min.      : 2072  Min.      :0.0050
## 1st Qu.:    4.00  1st Qu.:30285  1st Qu.:0.6254
## Median :   16.00  Median :47484  Median :0.6963
## Mean      :   84.77  Mean      :43664  Mean      :0.6791
## 3rd Qu.:   51.00  3rd Qu.:56718  3rd Qu.:0.7626
## Max.      :9491.00  Max.      :78841  Max.      :1.0000
## mean_mins_between_tweets mean_tweet_length mean_retweets
## Min.      :   -15.7  Min.      :  8.50  Min.      :    1.000
## 1st Qu.:   1152.8  1st Qu.: 80.79  1st Qu.:    1.167
## Median :   3851.7  Median : 91.74  Median :    1.636
## Mean      :  14715.4  Mean      : 91.41  Mean      :    3.873
## 3rd Qu.:  10823.8  3rd Qu.:103.28  3rd Qu.:    2.424
## Max.      :1139015.0  Max.      :287.88  Max.      :1961.300
## reply_rate
## Min.      :0.0000
## 1st Qu.:0.1232
## Median :0.3137
## Mean      :0.3411
## 3rd Qu.:0.5279
## Max.      :1.0000
```

From the summary, we can see that there are a couple factor variables in the data set, `bot`, `default_profile`, `default_profile_image` and `geo_enabled`. Before exploring further, let's first tell R that those columns represent categorical variables.

```
twitter$bot = factor(twitter$bot)
twitter$default_profile = factor(twitter$default_profile)
twitter$default_profile_image = factor(twitter$default_profile_image)
twitter$geo_enabled = factor(twitter$geo_enabled)

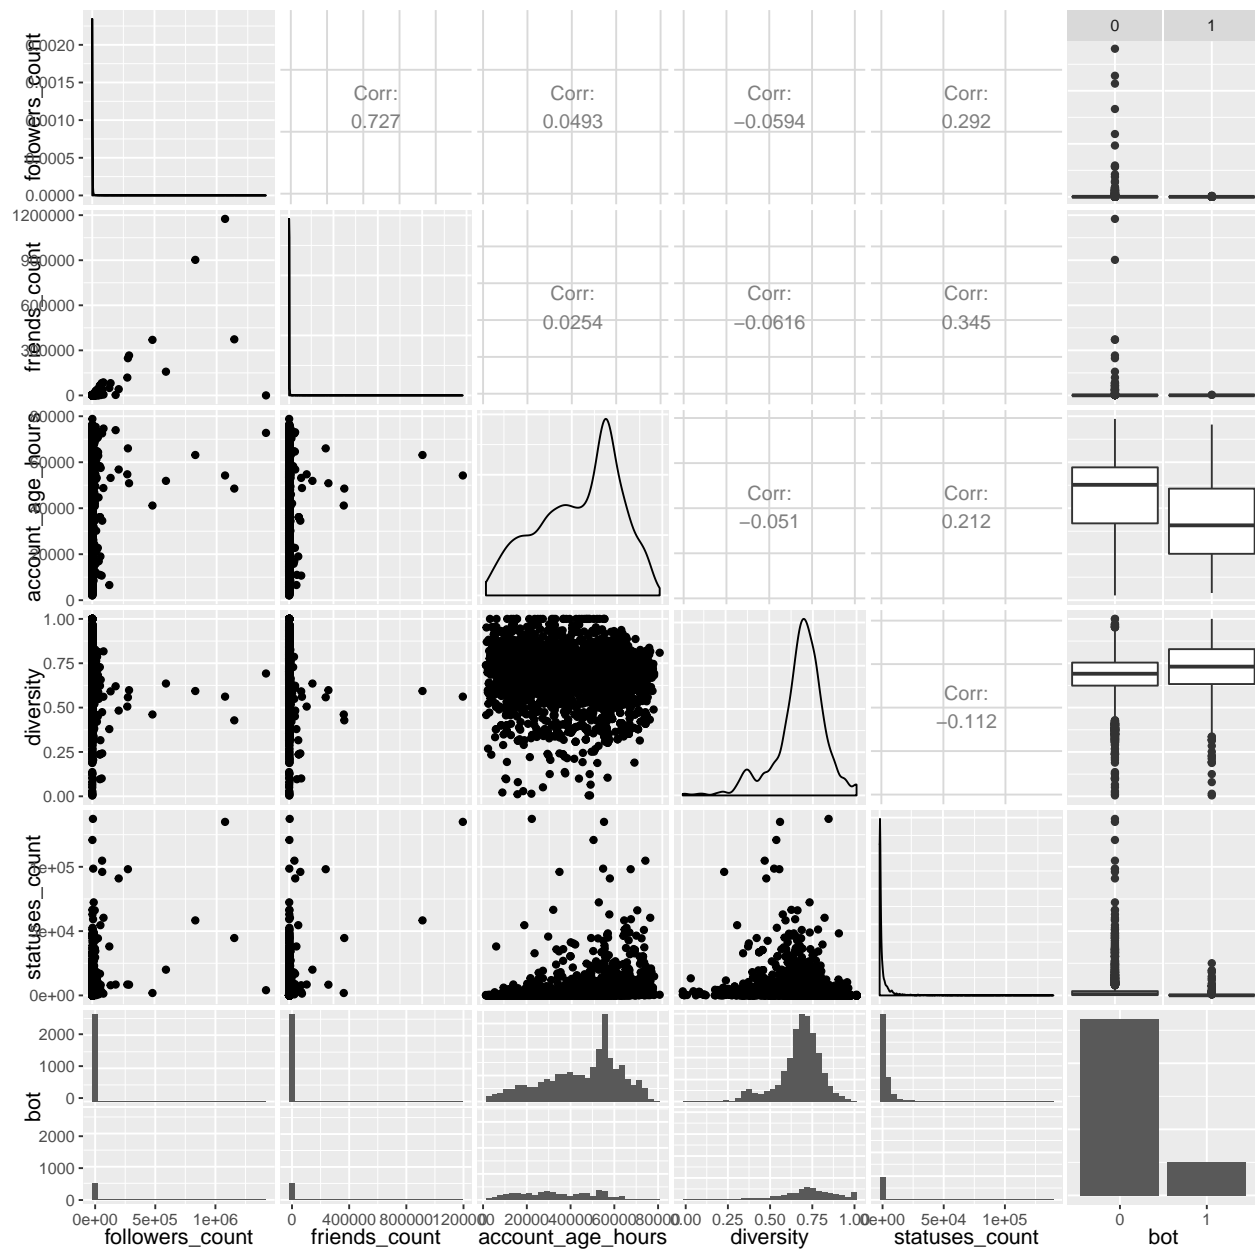
summary(twitter)
```

```
## bot      statuses_count      default_profile default_profile_image
## 0:2672  Min.      :    0  0:2256      0:3077
## 1: 504  1st Qu.:   188  1: 920      1: 99
##        Median :   723
##        Mean      :  3277
##        3rd Qu.:   2646
##        Max.      :137264
## friends_count      followers_count      favourites_count  geo_enabled
## Min.      :    11  Min.      :    0.0  Min.      :    0  0:1773
## 1st Qu.:    300  1st Qu.:    95.0  1st Qu.:    14  1:1403
## Median :    615  Median :   288.0  Median :    122
## Mean      :   2358  Mean      :  3709.3  Mean      :   1100
## 3rd Qu.:   1229  3rd Qu.:   830.5  3rd Qu.:    593
```

```
## Max. :1175187 Max. :1396699.0 Max. :176219
## listed_count account_age_hours diversity
## Min. : 0.00 Min. : 2072 Min. :0.0050
## 1st Qu.: 4.00 1st Qu.:30285 1st Qu.:0.6254
## Median : 16.00 Median :47484 Median :0.6963
## Mean : 84.77 Mean :43664 Mean :0.6791
## 3rd Qu.: 51.00 3rd Qu.:56718 3rd Qu.:0.7626
## Max. :9491.00 Max. :78841 Max. :1.0000
## mean_mins_between_tweets mean_tweet_length mean_retweets
## Min. : -15.7 Min. : 8.50 Min. : 1.000
## 1st Qu.: 1152.8 1st Qu.: 80.79 1st Qu.: 1.167
## Median : 3851.7 Median : 91.74 Median : 1.636
## Mean : 14715.4 Mean : 91.41 Mean : 3.873
## 3rd Qu.: 10823.8 3rd Qu.:103.28 3rd Qu.: 2.424
## Max. :1139015.0 Max. :287.88 Max. :1961.300
## reply_rate
## Min. :0.0000
## 1st Qu.:0.1232
## Median :0.3137
## Mean :0.3411
## 3rd Qu.:0.5279
## Max. :1.0000
```


Like before, we can evaluate many relationships simultaneously with `ggpairs`.

```
# inspect many trends with ggpairs
ggpairs(twitter[, c('followers_count', 'friends_count', 'account_age_hours',
                    'diversity', 'statuses_count', 'bot')])
```




Once we have some initial hypotheses we can make more specific plots.


```
ggplot(twitter, aes(x = followers_count, fill = bot)) +
  geom_histogram()
```




```
# Some people have a lot of followers, but most don't. we need to lob off  
# the long tail so we can see the distribution better  
ggplot(filter(twitter, followers_count < 100),  
  aes(x = followers_count, fill = bot)) +  
  geom_histogram()
```




```
ggplot(filter(twitter, followers_count < 100),  
  aes(x = followers_count, fill = bot)) +  
  geom_histogram() +  
  facet_grid(bot ~.)
```




```
# how about the number of people they follow?  
ggplot(twitter, aes(x = friends_count, fill = bot)) +  
  geom_histogram() +  
  facet_grid(bot ~.)
```




```
# it's a little hard to see  
ggplot(filter(twitter, friends_count < 2500),  
  aes(x = friends_count, fill = bot)) +  
  geom_histogram() +  
  facet_grid(bot ~.)
```

```
# what about account age?  
ggplot(twitter, aes(x = account_age_hours, fill = bot)) +  
  geom_histogram() +  
  facet_grid(bot ~.)
```




```
# lexical diversity
ggplot(twitter, aes(x = diversity, fill = bot)) +
  geom_histogram() +
  facet_grid(bot ~.)
```



```
# what are the average values?
avg_diversity =
  twitter %>%
    group_by(bot) %>%
    summarize(avg_diversity = mean(diversity, na.rm = TRUE))


# add it to the plot
ggplot(twitter, aes(x = diversity, fill = bot)) +
  geom_histogram() +
  geom_vline(data = avg_diversity, aes(xintercept = avg_diversity)) +
  facet_grid(bot ~.)
```



Feature engineering

Feature engineering is the process of creating predictor variables using domain knowledge. We can test hypotheses about the importance of various relationships by creating new predictors that help interrogate those relationships. For example, you might hypothesize a relationship between the number of tweets made and the lexical diversity that is relevant to model. To test that, make a new categorical variable indicating whether an account holder is a 'heavy tweeter', 'medium tweeter' or 'light tweeter':

```
# the number of tweets per account has a long tail  
ggplot(twitter, aes(x = statuses_count)) +  
  geom_histogram()
```



```
# break into three categories by quantile
quantile(twitter$statuses_count)
```

```
##      0%      25%      50%      75%     100%
##      0.0    188.0    723.0   2646.5 137264.0
```


```
# low tweeters will be the bottom 25%,
twitter$tweet_volume = NA
twitter$tweet_volume = ifelse(twitter$statuses_count <= 188,
                              'Light Tweeter',
                              twitter$tweet_volume)

twitter$tweet_volume = ifelse((twitter$statuses_count > 188 & twitter$statuses_count <= 2646),
                              'Medium Tweeter',
                              twitter$tweet_volume)

twitter$tweet_volume = ifelse(twitter$statuses_count > 2646,
                              'Heavy Tweeter',
                              twitter$tweet_volume)

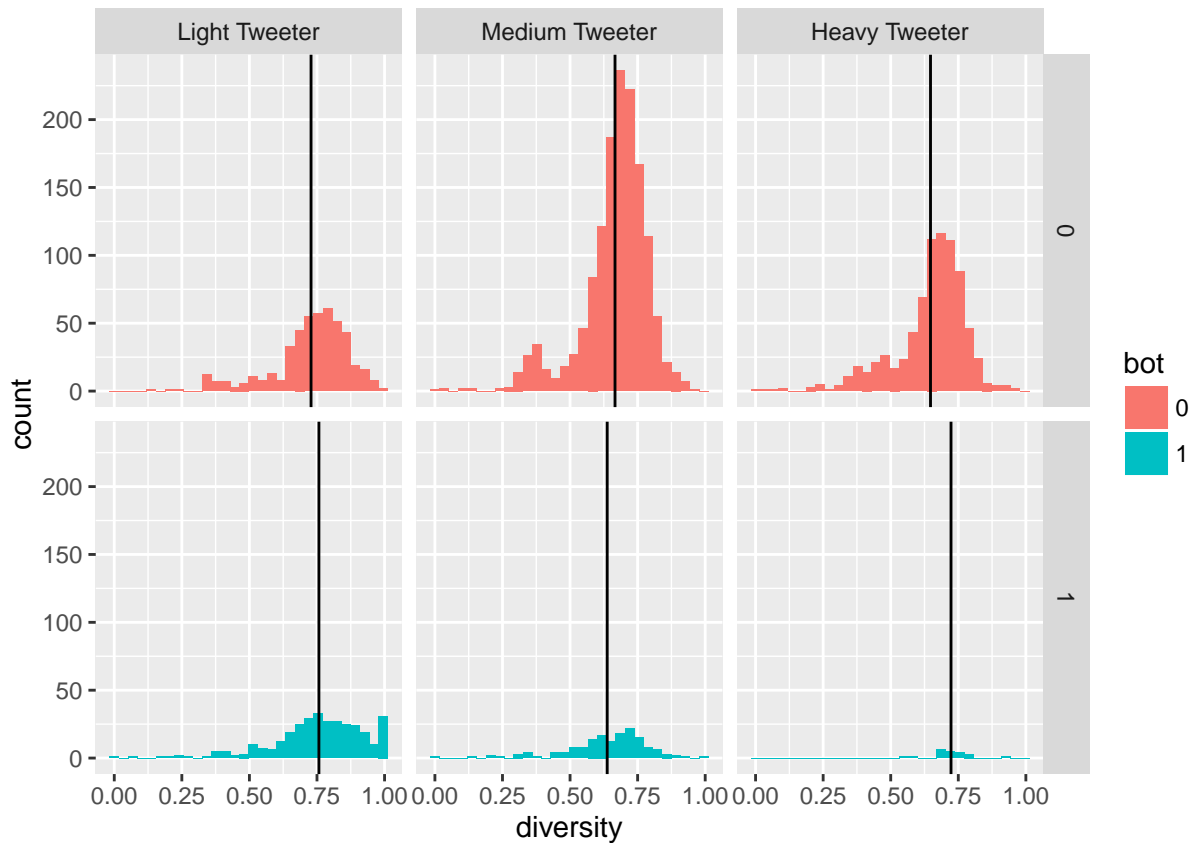
twitter$tweet_volume = factor(twitter$tweet_volume, levels = c('Light Tweeter', 'Medium Tweeter', 'Heavy Tweeter'))

# plot it!
ggplot(twitter, aes(x = statuses_count)) +
  geom_histogram(aes(fill = tweet_volume), bins = 100)
```



```
# update the figure
avg_diversity =
  twitter %>%
    group_by(bot, tweet_volume) %>%
    summarize(avg_diversity = mean(diversity, na.rm = TRUE))

ggplot(twitter, aes(x = diversity, fill = bot)) +
  geom_histogram() +
  geom_vline(data = avg_diversity, aes(xintercept = avg_diversity)) +
  facet_grid(bot ~ tweet_volume)
```



Logistic Regression

Training and testing sets

```
set.seed(243)
twitter = na.omit(twitter)

# select the training observations
in_train = createDataPartition(y = twitter$bot,
                                p = 0.75, # 75% in train, 25% in test
                                list = FALSE)

train = twitter[in_train, ]
test = twitter[-in_train, ]
```

Training logistic regressions

Check out this page for more types of logistic regression to try out.

```
logistic_model = train(bot ~ .,
                        data = train,
                        method = 'glm',
```

```

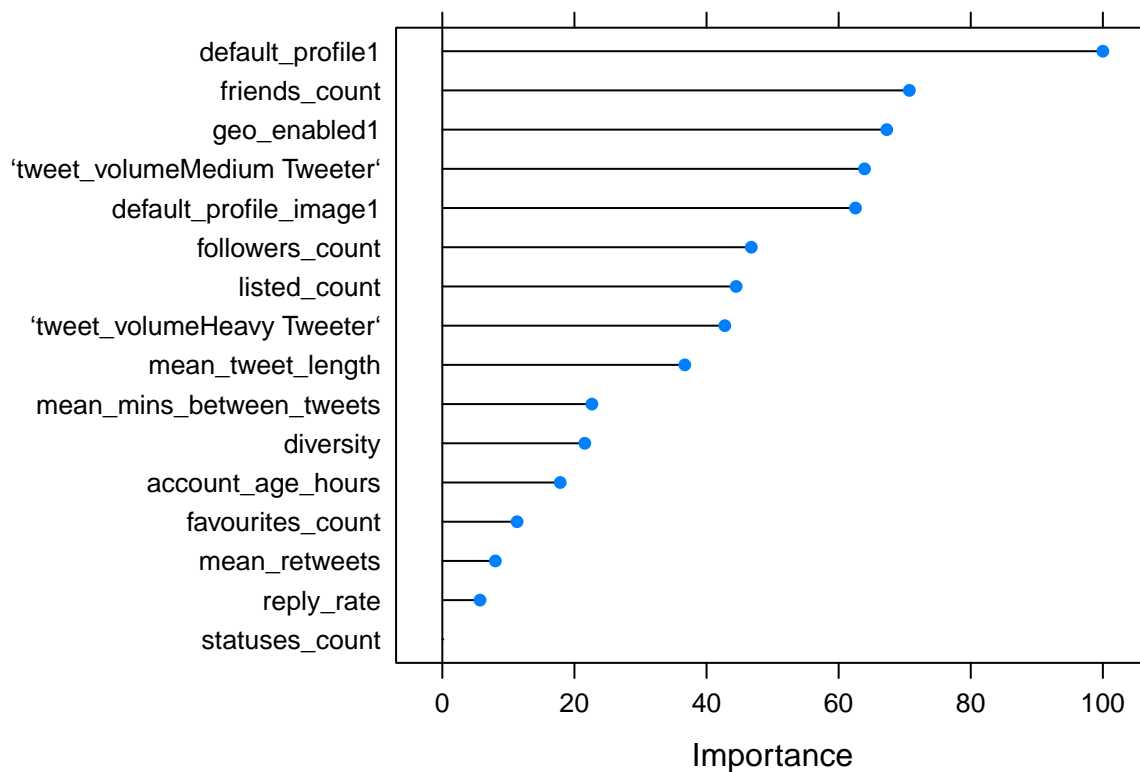
family = binomial,
preProcess = c('center', 'scale'))

summary(logistic_model)

##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8057  -0.4543  -0.2644  -0.1400   4.2037
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.66773    0.51971  -7.057 1.70e-12 ***
## statuses_count -0.05628    0.29941  -0.188 0.850892
## default_profile1  0.58481    0.06572   8.898 < 2e-16 ***
## default_profile_image1  0.52817    0.09371   5.636 1.74e-08 ***
## friends_count   23.55816    3.71218   6.346 2.21e-10 ***
## followers_count -40.02490    9.39093  -4.262 2.03e-05 ***
## favourites_count -0.45419    0.38694  -1.174 0.240474
## geo_enabled1    -0.51964    0.08592  -6.048 1.47e-09 ***
## listed_count     3.25166    0.80043   4.062 4.86e-05 ***
## account_age_hours -0.13779    0.07907  -1.743 0.081399 .
## diversity        0.14801    0.07161   2.067 0.038742 *
## mean_mins_between_tweets 0.18157    0.08404   2.161 0.030729 *
## mean_tweet_length -0.23152    0.06839  -3.385 0.000711 ***
## mean_retweets    -1.07502    1.21159  -0.887 0.374925
## reply_rate       -0.05054    0.07376  -0.685 0.493177
## `tweet_volumeMedium Tweeter` -0.47440    0.08243  -5.756 8.64e-09 ***
## `tweet_volumeHeavy Tweeter` -0.66668    0.17038  -3.913 9.12e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2084.2  on 2381  degrees of freedom
## Residual deviance: 1358.7  on 2365  degrees of freedom
## AIC: 1392.7
##
## Number of Fisher Scoring iterations: 15

plot(varImp(logistic_model))

```

```
# test predictions
```

```
logistic_predictions = predict(logistic_model, newdata = test)
confusionMatrix(logistic_predictions, test$bot)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  0    1
```

```
##           0 660  81
```

```
##           1   8  45
```

```
##
```

```
##           Accuracy : 0.8879
```

```
##           95% CI : (0.8639, 0.909)
```

```
##           No Information Rate : 0.8413
```

```
##           P-Value [Acc > NIR] : 0.0001104
```

```
##
```

```
##           Kappa : 0.4512
```

```
##           McNemar's Test P-Value : 2.312e-14
```

```
##
```

```
##           Sensitivity : 0.9880
```

```
##           Specificity : 0.3571
```

```
##           Pos Pred Value : 0.8907
```

```
##           Neg Pred Value : 0.8491
```

```
##           Prevalence : 0.8413
```

```
##           Detection Rate : 0.8312
```

```
##           Detection Prevalence : 0.9332
```

```
##           Balanced Accuracy : 0.6726
```

```
##
```

```
##      'Positive' Class : 0
##
```

There are subset selection methods for logistic regression as well. Try out `method = 'glmStepAIC'`:

```
# stepwise logisitic regression
step_model = train(bot ~ .,
                   data = train,
                   method = 'glmStepAIC',
                   family = binomial,
                   preProcess = c('center', 'scale'))
```

```
summary(step_model)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8225  -0.4520  -0.2657  -0.1406   4.1964
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.67541    0.51856  -7.088 1.36e-12 ***
## default_profile1    0.58261    0.06562   8.879 < 2e-16 ***
## default_profile_image1  0.52775    0.09347   5.646 1.64e-08 ***
## friends_count    23.78836    3.70813   6.415 1.41e-10 ***
## followers_count  -40.30223    9.35289  -4.309 1.64e-05 ***
## favourites_count  -0.50388    0.37668  -1.338 0.180997
## geo_enabled1     -0.51986    0.08588  -6.053 1.42e-09 ***
## listed_count      3.25599    0.80530   4.043 5.27e-05 ***
## account_age_hours -0.14007    0.07901  -1.773 0.076256 .
## diversity         0.14543    0.07162   2.031 0.042300 *
## mean_mins_between_tweets  0.18105    0.08431   2.147 0.031762 *
## mean_tweet_length  -0.22615    0.06810  -3.321 0.000897 ***
## mean_retweets     -1.14154    1.21830  -0.937 0.348761
## `tweet_volumeMedium Tweeter` -0.48019    0.08206  -5.851 4.87e-09 ***
## `tweet_volumeHeavy Tweeter` -0.69051    0.15135  -4.562 5.06e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2084.2  on 2381  degrees of freedom
## Residual deviance: 1359.2  on 2367  degrees of freedom
## AIC: 1389.2
##
## Number of Fisher Scoring iterations: 15
```

```
step_predictions = predict(step_model, newdata = test)
confusionMatrix(step_predictions, test$bot)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 660  79
##           1   8  47
##
##           Accuracy : 0.8904
##           95% CI : (0.8666, 0.9113)
##           No Information Rate : 0.8413
##           P-Value [Acc > NIR] : 4.658e-05
##
##           Kappa : 0.468
##           Mcnemar's Test P-Value : 6.153e-14
##
##           Sensitivity : 0.9880
##           Specificity : 0.3730
##           Pos Pred Value : 0.8931
##           Neg Pred Value : 0.8545
##           Prevalence : 0.8413
##           Detection Rate : 0.8312
##           Detection Prevalence : 0.9307
##           Balanced Accuracy : 0.6805
##
##           'Positive' Class : 0
##
```

How do the models compare?

```
# compare
results = resamples(list(logistic_model = logistic_model,
                        step_model = step_model))

# compare accuracy and kappa
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: logistic_model, step_model
## Number of resamples: 25
##
## Accuracy
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## logistic_model 0.8002  0.8777 0.8832 0.8780  0.8897 0.8976    0
## step_model     0.8654  0.8753 0.8815 0.8842  0.8939 0.9036    0
##
## Kappa
##           Min. 1st Qu. Median   Mean 3rd Qu.   Max. NA's
## logistic_model 0.2075  0.4444 0.4747 0.4513  0.4939 0.5407    0
## step_model     0.3869  0.4416 0.4706 0.4699  0.5001 0.5573    0
```

```
# plot results  
dotplot(results)
```

